

ECE 60146 HW6 Report

Zhengxin Jiang
(jiang839@purdue.edu)

1 Images From My Own Dataset




Figure 1: Three images from the class of bus




Figure 2: Three images from the class of cat




Figure 3: Three images from the class of pizza

2 Implementation of The Deep Network

For the deep network in this homework, we use the framework and the skipblock from last homework. The output of the network has 1620 neurons, and can be shaped into `num_cells * num_anchor_boxes * length_of_yolo_vector = 36 * 5 * 9`. The set of cells and anchor boxes are the same as in lecture notes.

```
212 class HW6Net(nn.Module):
213
214     def __init__(self, ngf=8, n_blocks=4):
215         super(HW6Net, self).__init__()
216
217         # The first conv layer
218         model = [nn.ReflectionPad2d(3),
219                 nn.Conv2d(3, ngf, kernel_size=7, padding=0),
220                 nn.BatchNorm2d(ngf),
221                 nn.ReLU(True)]
222
223         # Add downsampling layers
224         n_downsampling = 5
225         for i in range(n_downsampling):
226             mult = 2**i
227             model += [nn.Conv2d(ngf*mult, ngf*mult*2, kernel_size=3, stride=2, padding=1),
228                     nn.BatchNorm2d(ngf*mult*2),
229                     nn.ReLU(True)]
230
231         # My own ResNet blocks
232         mult = 2**n_downsampling
233         for i in range(n_blocks):
234             model += [ResnetBlock(ngf*mult, ngf*mult)]
235         self.model = nn.Sequential(*model)
236
237         wh = int(256/mult)
238
239         # output a yolo tensor
240         yolo_head = [nn.Flatten(),
241                     nn.Linear(ngf*mult*wh*wh, 4096),
242                     nn.ReLU(True),
243                     nn.Linear(4096, 2048),
244                     nn.ReLU(True),
245                     nn.Linear(2048, 1620)]
246         self.yolo_head = nn.Sequential(*yolo_head)
247
248     def forward(self, x):
249         ft = self.model(x)
250         yolo_tensor = self.yolo_head(ft)
251
252         return yolo_tensor
```

Figure 4: My implementation of the deep network

3 Training Logic and Validation Results

3.1 Implementation of Training Logic

First, the implementation of the `__getitem__` method of my dataset returns only the image tensor and the ground truth yolo_tensor. The calculation of ground truth yolo_tensor is done in the dataset's `__getitem__` function.

In the training process, the indices of yolo_vectors with an object presented are first extracted by using the `torch.nonzero()` function. We can then easily index the yolo_vectors in the further batch processing. Also, the information of the cell and anchor box indices are contained in the extracted indices themselves, so we don't need anything except the ground truth yolo_tensor from our dataloader.

```
34     ## get indices of yolo vectors
35     obj_indices = torch.nonzero(yoloTensor[:, :, :, 0])
36     nonobj_indices = torch.nonzero(yoloTensor[:, :, :, 0]==0)
```

Figure 5: Use `torch.nonzero` to extract the indices of yolo vectors

For the validation part, we put both the predicted and ground truth boxes/labels on the validation images and see how the trained network perform.

3.2 Plot of Training Loss

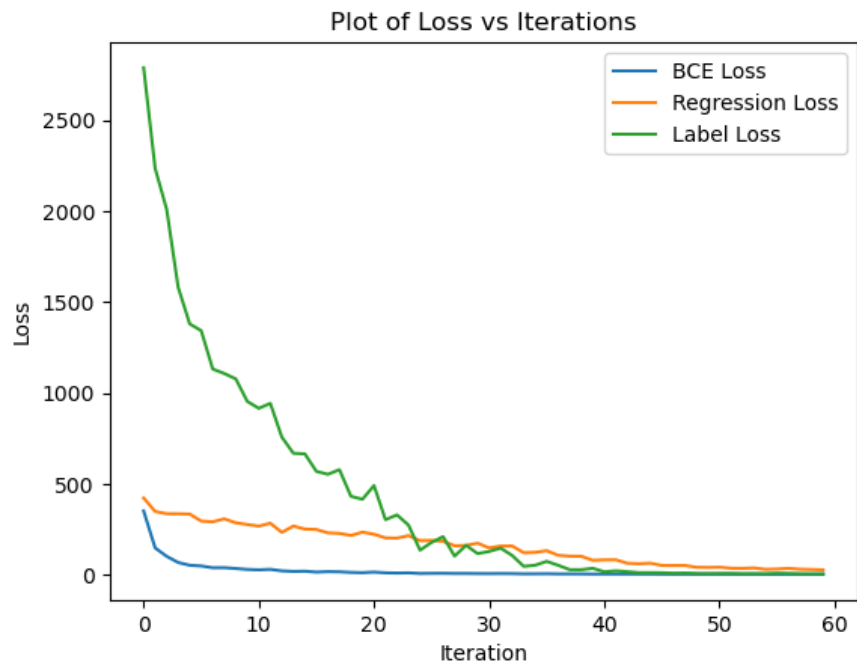


Figure 6: Plot of training loss vs iterations

3.3 Images of Validation Results



Figure 7: Three validation results from the class of bus




Figure 8: Three validation results from the class of cat




Figure 9: Three validation results from the class of pizza

4 Discussion

The network basically detects most objects correctly. One thing I observed is that the network tends to give a relatively small bbox when the ground truth bbox is large and nearly square-shaped. This might due to the limitation of the setting of anchor boxes.

5 Source code

```
# ECE60146 HW6
# Zhengxin Jiang
# jiang839

import numpy as np
import os
import matplotlib.pyplot as plt
from PIL import Image
from pycocotools.coco import COCO
import seaborn as sn
import random
import json
import math

import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision
import torchvision.transforms as tvt
from torch.utils.data import DataLoader

import cv2

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)

# Function for preparing the image data
def prepData(coco, root, tasktype):

    if tasktype == 'training':
        rawDataDir = os.path.join(root, 'train2014').replace("\\", "/")
        hwDataDir = os.path.join(root, 'hw6', 'train').replace("\\", "/")

    if tasktype == 'validation':
        rawDataDir = os.path.join(root, 'val2014').replace("\\", "/")
        hwDataDir = os.path.join(root, 'hw6', 'val').replace("\\", "/")

    bbox_total = []
    label_total = []
    selectedImgIds = []
    imgCount = 0

    catIds = coco.getCatIds(catNms=['bus', 'cat', 'pizza'])
    coco_labels_inverse = {}
    for i, catid in enumerate(catIds):
        coco_labels_inverse[catid] = i

    # go through the images
    for catCount, catId in enumerate(catIds):

        ImgIds = coco.getImgIds(catIds=catId)
```

```

for imgId in ImgIds:

    if imgId in selectedImgIds:
        continue

    coco_img = coco.loadImgs(imgId)[0]
    annId = coco.getAnnIds(imgIds=coco_img['id'], catIds=catIds, iscrowd=None)
    anns = coco.loadAnns(annId)

    # check if the image is valid
    is_valid = 0
    anns_valid = []

    for ann in anns:
        if ann['area'] > 4096:
            is_valid = 1
            anns_valid.append(ann)

    if is_valid==0:
        continue

    imgName = coco_img['file_name']
    img = Image.open(os.path.join(rawDataDir, imgName).replace("\\", "/"))

    # save bbox and labels
    img_bbox = []
    img_labels = []

    for ann in anns_valid:
        bbox = ann['bbox']
        label = coco_labels_inverse[ann['category_id']]

        resize_ratio = (256/img.size[0], 256/img.size[1])
        bbox_resized = np.zeros(4)
        bbox_resized[0] = bbox[0]*resize_ratio[0]
        bbox_resized[2] = bbox[2]*resize_ratio[0]
        bbox_resized[1] = bbox[1]*resize_ratio[1]
        bbox_resized[3] = bbox[3]*resize_ratio[1]

        bbox_resized[2:] += bbox_resized[:2] # change the format to [x1, y1, x2, y2]
        bbox_resized[2:] -= 1e-3
        bbox_resized[:2] += 1e-3 # let bbox reside in (0,1)

        img_bbox.append(list(bbox_resized))
        img_labels.append(label)

    bbox_total.append(img_bbox)
    label_total.append(img_labels)

    if img.mode != "RGB":
        img = img.convert(mode="RGB")
    img = img.resize((256, 256), Image.BOX)

```

```

        imgNewName = str(imgCount) + '.jpg'
        fp = open(os.path.join(hwDataDir, imgNewName).replace("\\", "/"), 'w')
        img.save(fp)

        #update
        imgCount += 1
        selectedImgIds.append(imgId)

with open('{}/bbox'.format(hwDataDir), 'w') as fp:
    json.dump(bbox_total, fp)
with open('{}/labels'.format(hwDataDir), 'w') as fp:
    json.dump(label_total, fp)

return

# The Dataset class for hw5
class hwDataset(torch.utils.data.Dataset):

    def __init__(self, root, tasktype):
        super().__init__()

        if tasktype == 'training':
            self.root = os.path.join(root, 'hw6', 'train').replace("\\", "/")
        if tasktype == 'validation':
            self.root = os.path.join(root, 'hw6', 'val').replace("\\", "/")

        with open(os.path.join(self.root, 'bbox').replace("\\", "/"), 'r') as fp:
            self.bbox = json.load(fp)
        with open(os.path.join(self.root, 'labels').replace("\\", "/"), 'r') as fp:
            self.labels = json.load(fp)

        len_list = []
        for i, item in enumerate(self.labels):
            len_list.append(len(item))

        self.maxobjcount = max(len_list)

    def __len__(self):
        return len(os.listdir(self.root))-2 # Total files minus the two json file

    def __getitem__(self, index):

        name = str(index)+' .jpg'

        img = Image.open(os.path.join(self.root, name).replace("\\", "/"))

        tr = tvn.Compose([
            tvn.ToTensor(),
            tvn.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
        ])

```

```

img_tensor = tr(img)
label = self.labels[index]
bb_scaled = np.array(self.bbox[index])/256

cellIdx = []
aboxIdx = []
yoloVector = []

yolo_tensor = torch.zeros(num_yolo_cells, num_anchor_boxes, 8)

for i in range(len(label)):
    bb_w = bb_scaled[i][2] - bb_scaled[i][0]
    bb_h = bb_scaled[i][3] - bb_scaled[i][1]
    bb_center_w = (bb_scaled[i][0]+bb_scaled[i][2])/2
    bb_center_h = (bb_scaled[i][1]+bb_scaled[i][3])/2

    # 6*6 cells and 5 anchorbox for each cell
    cell_w = 1/6
    cellidx_i = int(bb_center_w / cell_w)
    cellidx_j = int(bb_center_h / cell_w)
    cellidx = cellidx_j*6 + cellidx_i
    cell_center_w = (cellidx_i + 0.5) * cell_w
    cell_center_h = (cellidx_j + 0.5) * cell_w
    del_w = bb_center_w - cell_center_w
    del_h = bb_center_h - cell_center_h

    ar = bb_h / bb_w
    if ar <= 0.2:
        aboxidx = 0
        sigma_w = math.log(bb_w/(cell_w*5))
        sigma_h = math.log(bb_h/(cell_w))
    if 0.2 < ar <= 0.5:
        aboxidx = 1
        sigma_w = math.log(bb_w/(cell_w*3))
        sigma_h = math.log(bb_h/(cell_w))
    if 0.5 < ar <= 1.5:
        aboxidx = 2
        sigma_w = math.log(bb_w/(cell_w))
        sigma_h = math.log(bb_h/(cell_w))
    if 1.5 < ar <= 4.0:
        aboxidx = 3
        sigma_w = math.log(bb_w/(cell_w))
        sigma_h = math.log(bb_h/(cell_w*3))
    if 4.0 < ar :
        aboxidx = 4
        sigma_w = math.log(bb_w/(cell_w))
        sigma_h = math.log(bb_h/(cell_w*5))

    # construct yolo vector
    yolo_vector = torch.tensor([1, del_w, del_h, sigma_w, sigma_h, 0, 0, 0])
    yolo_vector[5+label[i]] = 1

    yolo_tensor[cellidx, aboxidx] = yolo_vector

```



```

        return img_tensor, yolo_tensor

# The Resnet block with skip connection
class ResnetBlock(nn.Module):

    def __init__(self, in_ch, out_ch):
        super(ResnetBlock, self).__init__()

        self.conv = nn.Conv2d(in_ch, out_ch, kernel_size=3, stride=1, padding=1)
        self.bn = nn.BatchNorm2d(out_ch)

    def forward(self, x):
        identity = x

        out = F.relu(self.bn(self.conv(x)))
        out = self.bn(self.conv(out))
        out += identity #skip connection
        out = F.relu(out)

        return out

class HW6Net(nn.Module):

    def __init__(self, ngf=8, n_blocks=4):
        super(HW6Net, self).__init__()

        # The first conv layer
        model = [nn.ReflectionPad2d(3),
                  nn.Conv2d(3, ngf, kernel_size=7, padding=0),
                  nn.BatchNorm2d(ngf),
                  nn.ReLU(True)]

        # Add downsampling layers
        n_downsampling = 5
        for i in range(n_downsampling):
            mult = 2**i
            model += [nn.Conv2d(ngf*mult, ngf*mult*2, kernel_size=3, stride=2, padding=1),
                      nn.BatchNorm2d(ngf*mult*2),
                      nn.ReLU(True)]

        # My own ResNet blocks
        mult = 2**n_downsampling
        for i in range(n_blocks):
            model += [ResnetBlock(ngf*mult, ngf*mult)]
        self.model = nn.Sequential(*model)

        wh = int(256/mult)

        # output a yolo tensor
        yolo_head = [nn.Flatten(),

```

```

        nn.Linear(ngf*mult*wh*wh, 4096),
        nn.ReLU(True),
        nn.Linear(4096, 2048),
        nn.ReLU(True),
        nn.Linear(2048, 1620)]
    self.yolo_head = nn.Sequential(*yolo_head)

    def forward(self, x):
        ft = self.model(x)
        yolotensor = self.yolo_head(ft)

        return yolotensor

# Recover the original bbox
def bboxRecovery(cell_idx, anchor_idx, yolovec):

    cell_idx_i = cell_idx % 6
    cell_idx_j = cell_idx // 6

    cell_w = 1/6
    bbox_center_w = (cell_idx_i+0.5)*cell_w + yolovec[1]
    bbox_center_h = (cell_idx_j+0.5)*cell_w + yolovec[2]

    if anchor_idx == 0:
        bbox_w = math.exp(yolovec[3])*cell_w*5
        bbox_h = math.exp(yolovec[4])*cell_w
    if anchor_idx == 1:
        bbox_w = math.exp(yolovec[3])*cell_w*3
        bbox_h = math.exp(yolovec[4])*cell_w
    if anchor_idx == 2:
        bbox_w = math.exp(yolovec[3])*cell_w
        bbox_h = math.exp(yolovec[4])*cell_w
    if anchor_idx == 3:
        bbox_w = math.exp(yolovec[3])*cell_w
        bbox_h = math.exp(yolovec[4])*cell_w*3
    if anchor_idx == 4:
        bbox_w = math.exp(yolovec[3])*cell_w
        bbox_h = math.exp(yolovec[4])*cell_w*5

    bbox_raw = 256 * np.array([bbox_center_w - bbox_w/2, bbox_center_h - bbox_h/2,
        bbox_center_w + bbox_w/2, bbox_center_h + bbox_h/2])
    bbox_raw = bbox_raw.astype(int)

    return bbox_raw

def drawBBox(img, label, gt, color):

    class_name = ['bus', 'cat', 'pizza']

    if color == 'r':
        cl = (36, 12, 255)

```

```

    if color == 'g':
        c1 = (36, 255, 12)

    img = cv2.rectangle(img, (int(gt[0]), int(gt[1])), (int(gt[2]), int(gt[3])), c1, 2)
    img = cv2.putText(img, class_name[label], (int(gt[0]), int(gt[1]-10)), cv2.
        FONT_HERSHEY_SIMPLEX, 0.8, c1, 2)

    return img

# global variables
num_yolo_cells = 36
num_anchor_boxes = 5
batch_size = 4
epochs = 20

##### main #####
coco_train = COCO('{}train2014/annotations/instances_train2014.json'.format(root))
coco_val = COCO('{}val2014/annotations/instances_val2014.json'.format(root))

root = 'D:/coco'

prepData(coco_train, root, 'training')
prepData(coco_val, root, 'validation')

#### Training ####
traindataset = hwDataset(root, 'training')
train_data_loader = DataLoader(traindataset, batch_size=batch_size, num_workers=0,
    shuffle=True)

net = HW6Net()
net = net.to(device)

criterion1 = torch.nn.BCELoss()
criterion2 = torch.nn.MSELoss()
criterion3 = torch.nn.CrossEntropyLoss()

optimizer = torch.optim.SGD(net.parameters(), lr=1e-4, momentum=0.9)
# optimizer = torch.optim.Adam(net.parameters(), lr=1e-4, betas=(0.9, 0.99))

bce_loss_list = []
mse_loss_list = []
label_loss_list = []

for epoch in range(epochs):
    # running_loss = 0.0
    bce_running_loss = 0.0
    mse_running_loss = 0.0
    label_running_loss = 0.0

    for i, data in enumerate(train_data_loader):
        inputs, yoloTensor = data

```

```

inputs = inputs.to(device)
yoloTensor = yoloTensor.to(device)

yoloTensor_aug = torch.zeros(batch_size, num_yolo_cells, num_anchor_boxes, 9).to(
    device)
yoloTensor_aug[:, :, :, :-1] = yoloTensor

## get indices of yolo vectors
obj_indices = torch.nonzero(yoloTensor[:, :, :, 0])
nonobj_indices = torch.nonzero(yoloTensor[:, :, :, 0]==0)

## assign prob mass for no object
for nonobj_idx in nonobj_indices:
    yoloTensor_aug[tuple(nonobj_idx)][-1] = 1

optimizer.zero_grad()
outputs = net(inputs)
prediction_aug = outputs.view(batch_size, num_yolo_cells, num_anchor_boxes, 9)

# loss = torch.tensor(0.0, requires_grad=True).float().to(device)
BCEloss = torch.tensor(0.0, requires_grad=True).float().to(device)
MSEloss = torch.tensor(0.0, requires_grad=True).float().to(device)
Labelloss = torch.tensor(0.0, requires_grad=True).float().to(device)

for obj_idx in obj_indices:
    yolovec_pred = prediction_aug[tuple(obj_idx)]
    yolovec_gt = yoloTensor_aug[tuple(obj_idx)]

    object_presence = nn.Sigmoid()(torch.unsqueeze(yolovec_pred[0], dim=0))
    target = torch.unsqueeze(yolovec_gt[0], dim=0)
    bceloss = criterion1(object_presence, target)
    BCEloss += bceloss
# loss += bceloss

    pred_regression = torch.unsqueeze(yolovec_pred[1:5], dim=0)
    target_regression = torch.unsqueeze(yolovec_gt[1:5], dim=0)
    regloss = criterion2(pred_regression, target_regression)
    MSEloss += regloss
# loss += regloss

    pred_probvec = torch.unsqueeze(yolovec_pred[5:], dim=0)
    target_probvec = torch.unsqueeze(yolovec_gt[5:], dim=0)
    labelloss = criterion3(pred_probvec, target_probvec)
    Labelloss += labelloss
# loss += labelloss

BCEloss.backward(retain_graph=True)
MSEloss.backward(retain_graph=True)
Labelloss.backward()
# loss.backward()
optimizer.step()

```

```

        bce_running_loss += BCEloss.item()
        mse_running_loss += MSEloss.item()
        label_running_loss += Labelloss.item()
# running_loss += loss.item()
    if (i+1) % 500 == 0:
        print("[epoch: %d, batch: %5d] BCEloss: %.4f" % (epoch + 1, i + 1,
            bce_running_loss / 500))
        bce_loss_list.append(bce_running_loss)
        bce_running_loss = 0.0
        print("[epoch: %d, batch: %5d] MSEloss: %.4f" % (epoch + 1, i + 1,
            mse_running_loss / 500))
        mse_loss_list.append(mse_running_loss)
        mse_running_loss = 0.0
        print("[epoch: %d, batch: %5d] Labelloss: %.4f" % (epoch + 1, i + 1,
            label_running_loss / 500))
        label_loss_list.append(label_running_loss)
        label_running_loss = 0.0

torch.save(net.state_dict(), '20e.pth')

## plot the training loss
plt.plot(bce_loss_list, label='BCE Loss')
plt.plot(mse_loss_list, label='Regression Loss')
plt.plot(label_loss_list, label='Label Loss')
plt.legend()
plt.title("Plot of Loss vs Iterations")
plt.xlabel('Iteration')
plt.ylabel('Loss')
plt.show()

#### Validation ####
valdataset = hwDataset(root, 'validation')
val_data_loader = DataLoader(valdataset, batch_size=1, num_workers=0, shuffle=False)

# net = HW6Net()
# net = net.to(device)
# ## load trained parameters
# net.load_state_dict(torch.load('20e.pth', map_location=torch.device(device)))

with torch.no_grad():
    for i, data in enumerate(val_data_loader):

        inputs, yoloTensor = data

        inputs = inputs.to(device)

        outputs = net(inputs)
        prediction_aug = outputs.view(1, num_yolo_cells, num_anchor_boxes, 9)
        prediction_aug = prediction_aug.to('cpu')

        imgname = '{}.jpg'.format(i)
        img = cv2.imread(os.path.join(root, 'hw6', 'val', imgname).replace("\\", "/"))

```

```

## draw ground truth and prediction
obj_indices = torch.nonzero(yoloTensor[:, :, :, 0])
for obj_idx in obj_indices:
    cell_idx = obj_idx[1]
    anchbox_idx = obj_idx[2]
    yolovec_gt = yoloTensor[tuple(obj_idx)]
    yolovec_pred = prediction_aug[tuple(obj_idx)]

    bbox_gt_raw = bboxRecovery(cell_idx.item(), anchbox_idx.item(), yolovec_gt)
    label_gt = torch.max(yolovec_gt[5:8], 0)[1].item()
    img = drawBBox(img, label_gt, bbox_gt_raw, 'g')

    bbox_pred_raw = bboxRecovery(cell_idx.item(), anchbox_idx.item(), yolovec_pred
    )
    label_pred = torch.max(yolovec_pred[5:8], 0)[1].item()
    img = drawBBox(img, label_pred, bbox_pred_raw, 'r')

cv2.imwrite(os.path.join(root, 'hw6', 'valresult', imgname).replace("\\", "/"), img
)

```